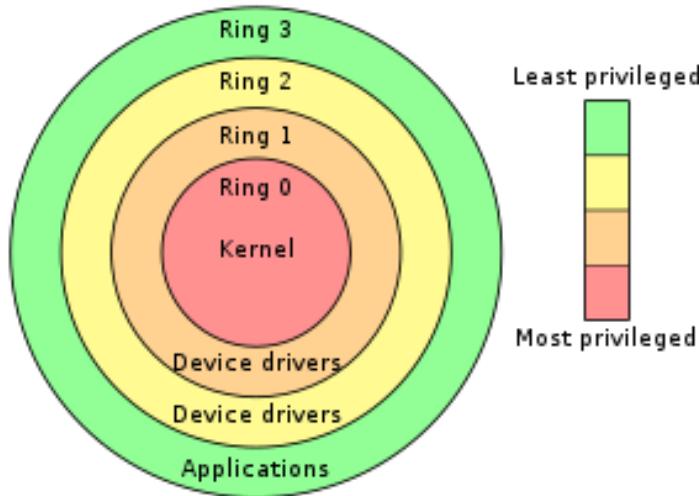


리눅스 커널 취약점 탐지 및 공격

Byoungyoung Lee (이병영)
Seoul National University (SNU)
byoungyoung@snu.ac.kr
<https://compsec.snu.ac.kr>

Kernel Vulnerability



https://en.wikipedia.org/wiki/Principle_of_least_privilege

**Exploiting kernel vulnerability, attacker can subvert the entire system.
Identifying and patching kernel vulnerabilities is very important!**

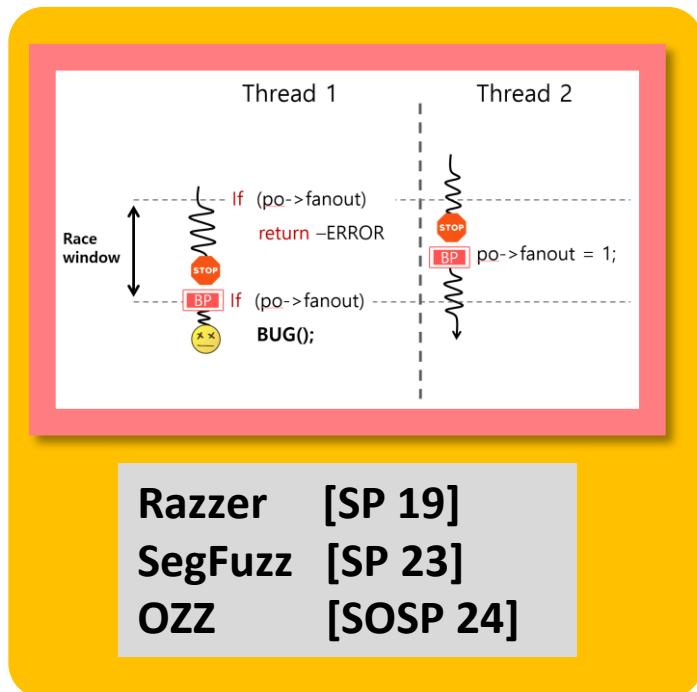
Linux Kernel and Security

- **Linux kernel is the backbone of modern computing infrastructures**
 - Desktop
 - Cloud host and guest VM
 - Android
 - Various IoT devices
 - Many more...

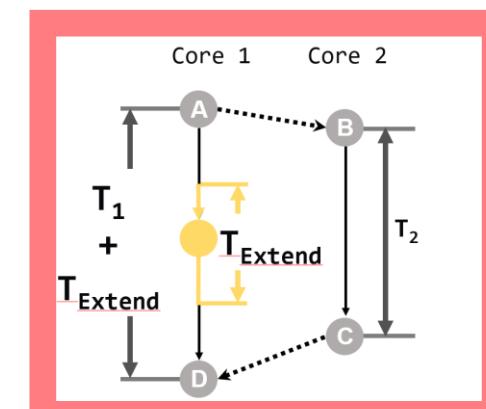
Examples of Linux Kernel Vulnerabilities



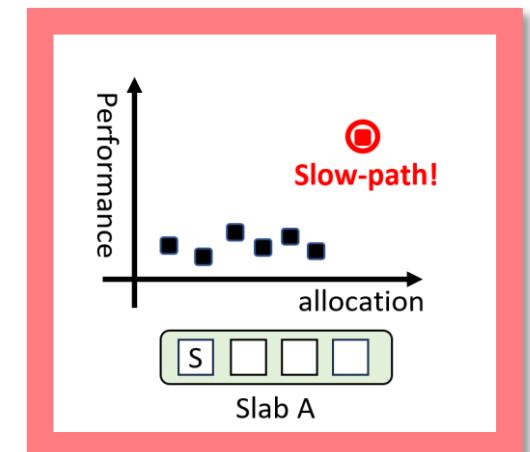
Our Research on Linux Kernel Security



Fuzzing to find race vulnerabilities



ExpRace [Security 21]



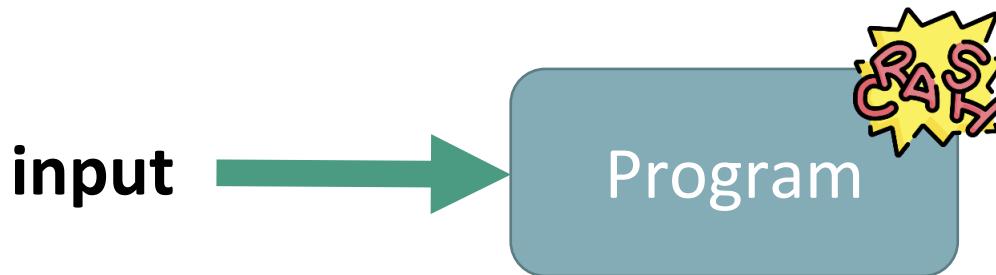
Pspray [Security 23]

Exploiting use-after-free

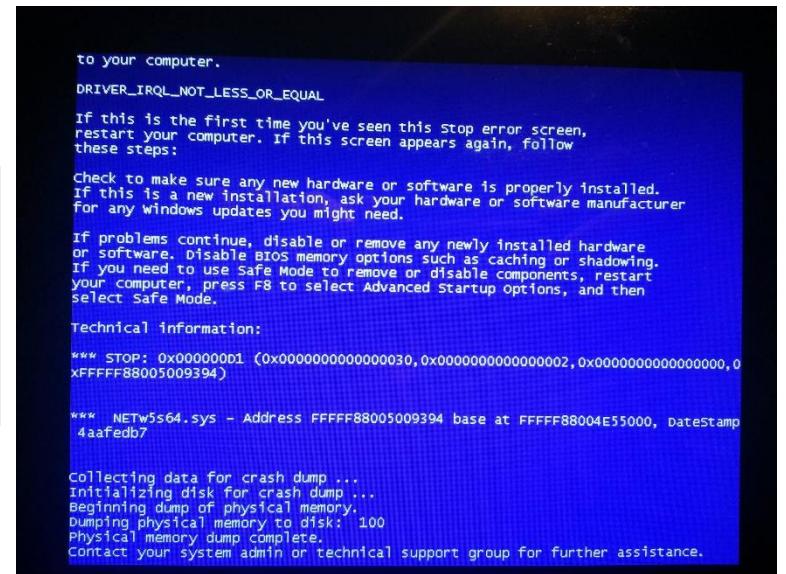
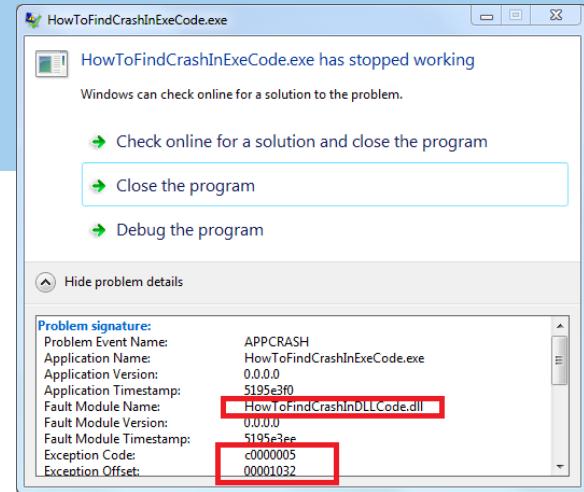
What programs do



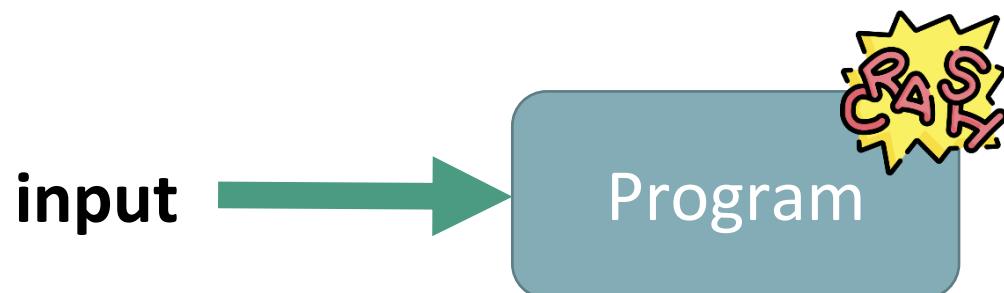
Vulnerability



For some input, the program acts weird
(e.g., crash, hanging, etc.)



Fuzzing: Random Testing!



Fuzzing 101

- #1. Randomly pick input $x^R \in X$
- #2. Run the program using x^R
 - If something bad occurs,
 x^R induces a vulnerability
- #3. Go back to step #1

Why is fuzzing so popular today?

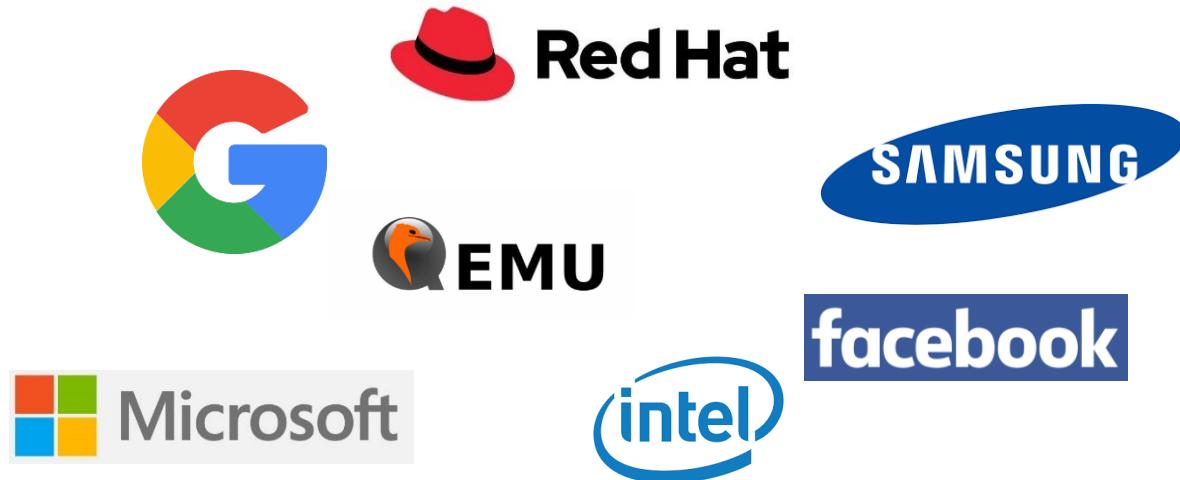


Easy-to-implement



Effective

Real-world Practices of Fuzzing



Almost all big-techs fuzz their products

Race condition

- **Race conditions**
 - If following three conditions meet
 - Two instructions access the same memory location
 - At least one of two is a write instruction
 - These two are executed concurrently
 - If a race condition occurs, the computational results may vary depending on the execution order.

Race condition: A simple example

```
a = 0;
```

```
If (a) {  
    b = 0;  
} else {  
    b = 1;  
}
```

Thread 1

b will be 1

```
a = 1;
```

```
a = 0;
```

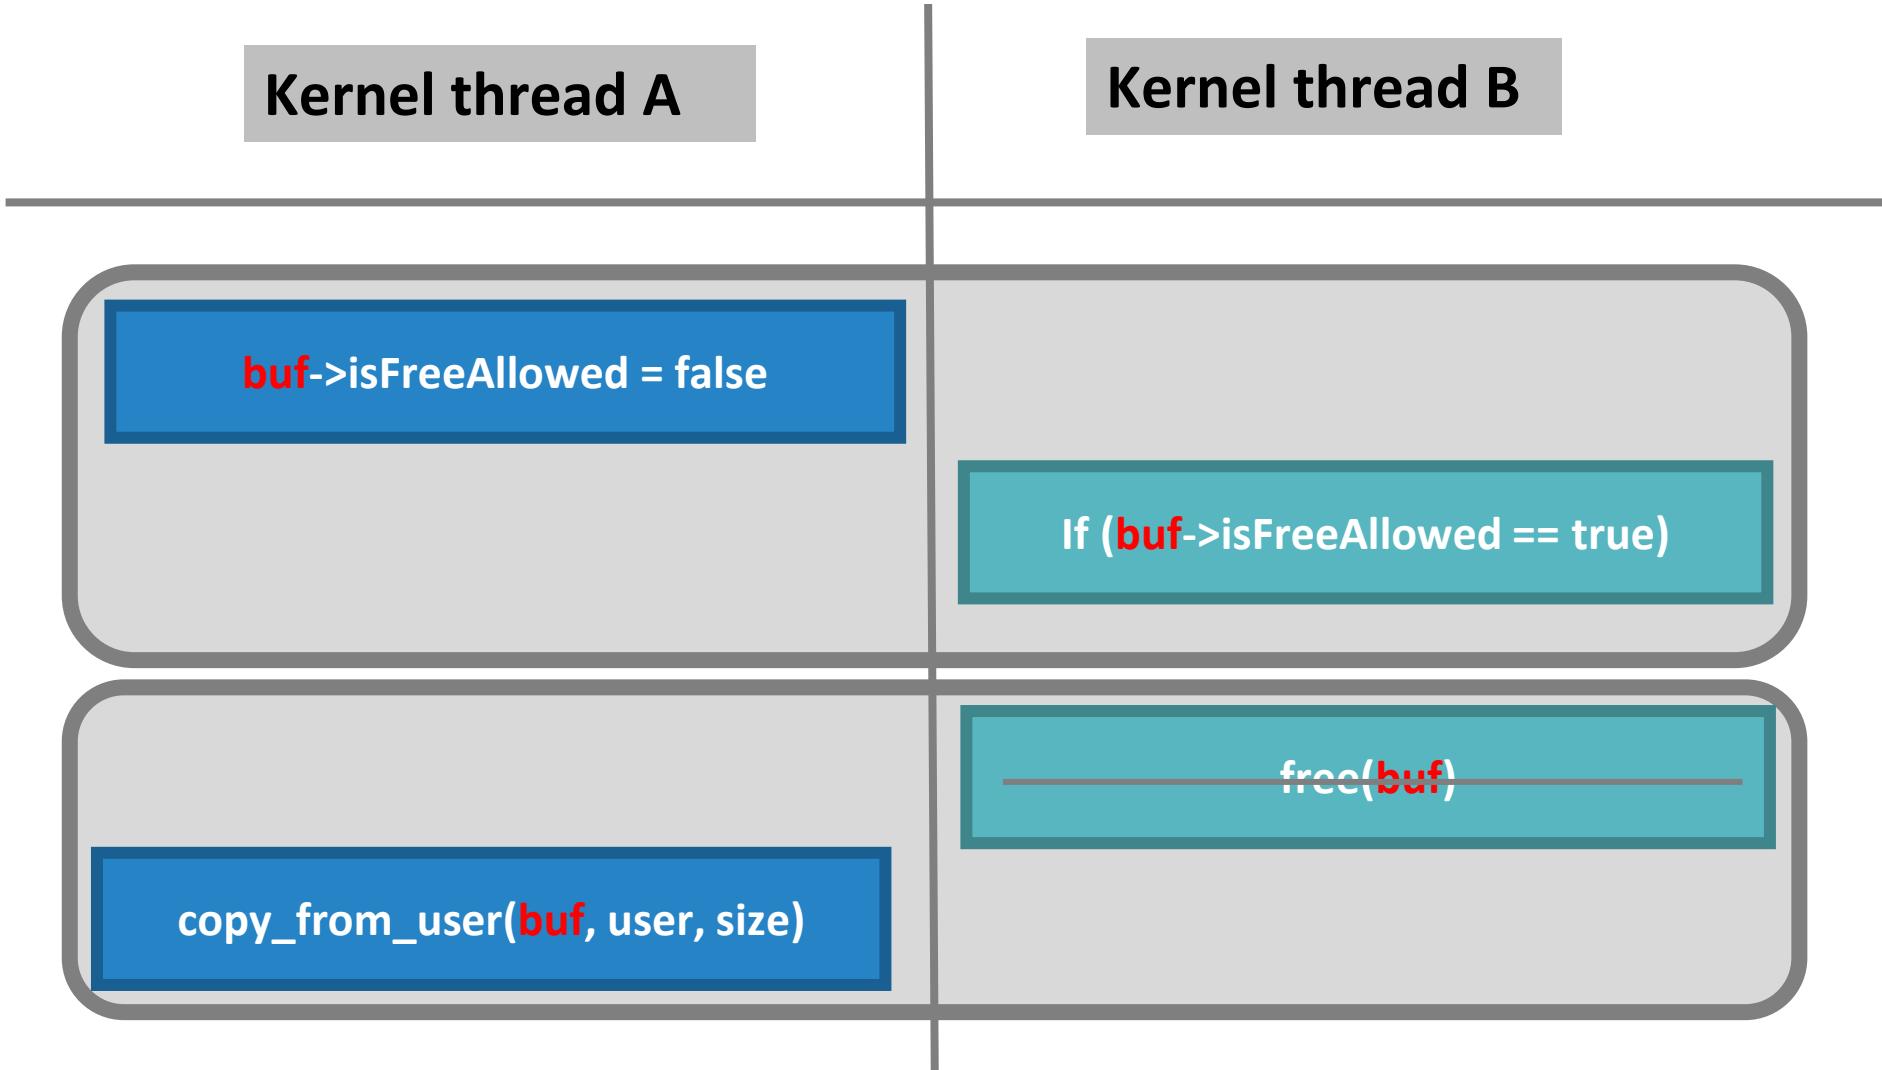
```
If (a) {  
    b = 0;  
} else {  
    b = 1;  
}
```

Thread 1

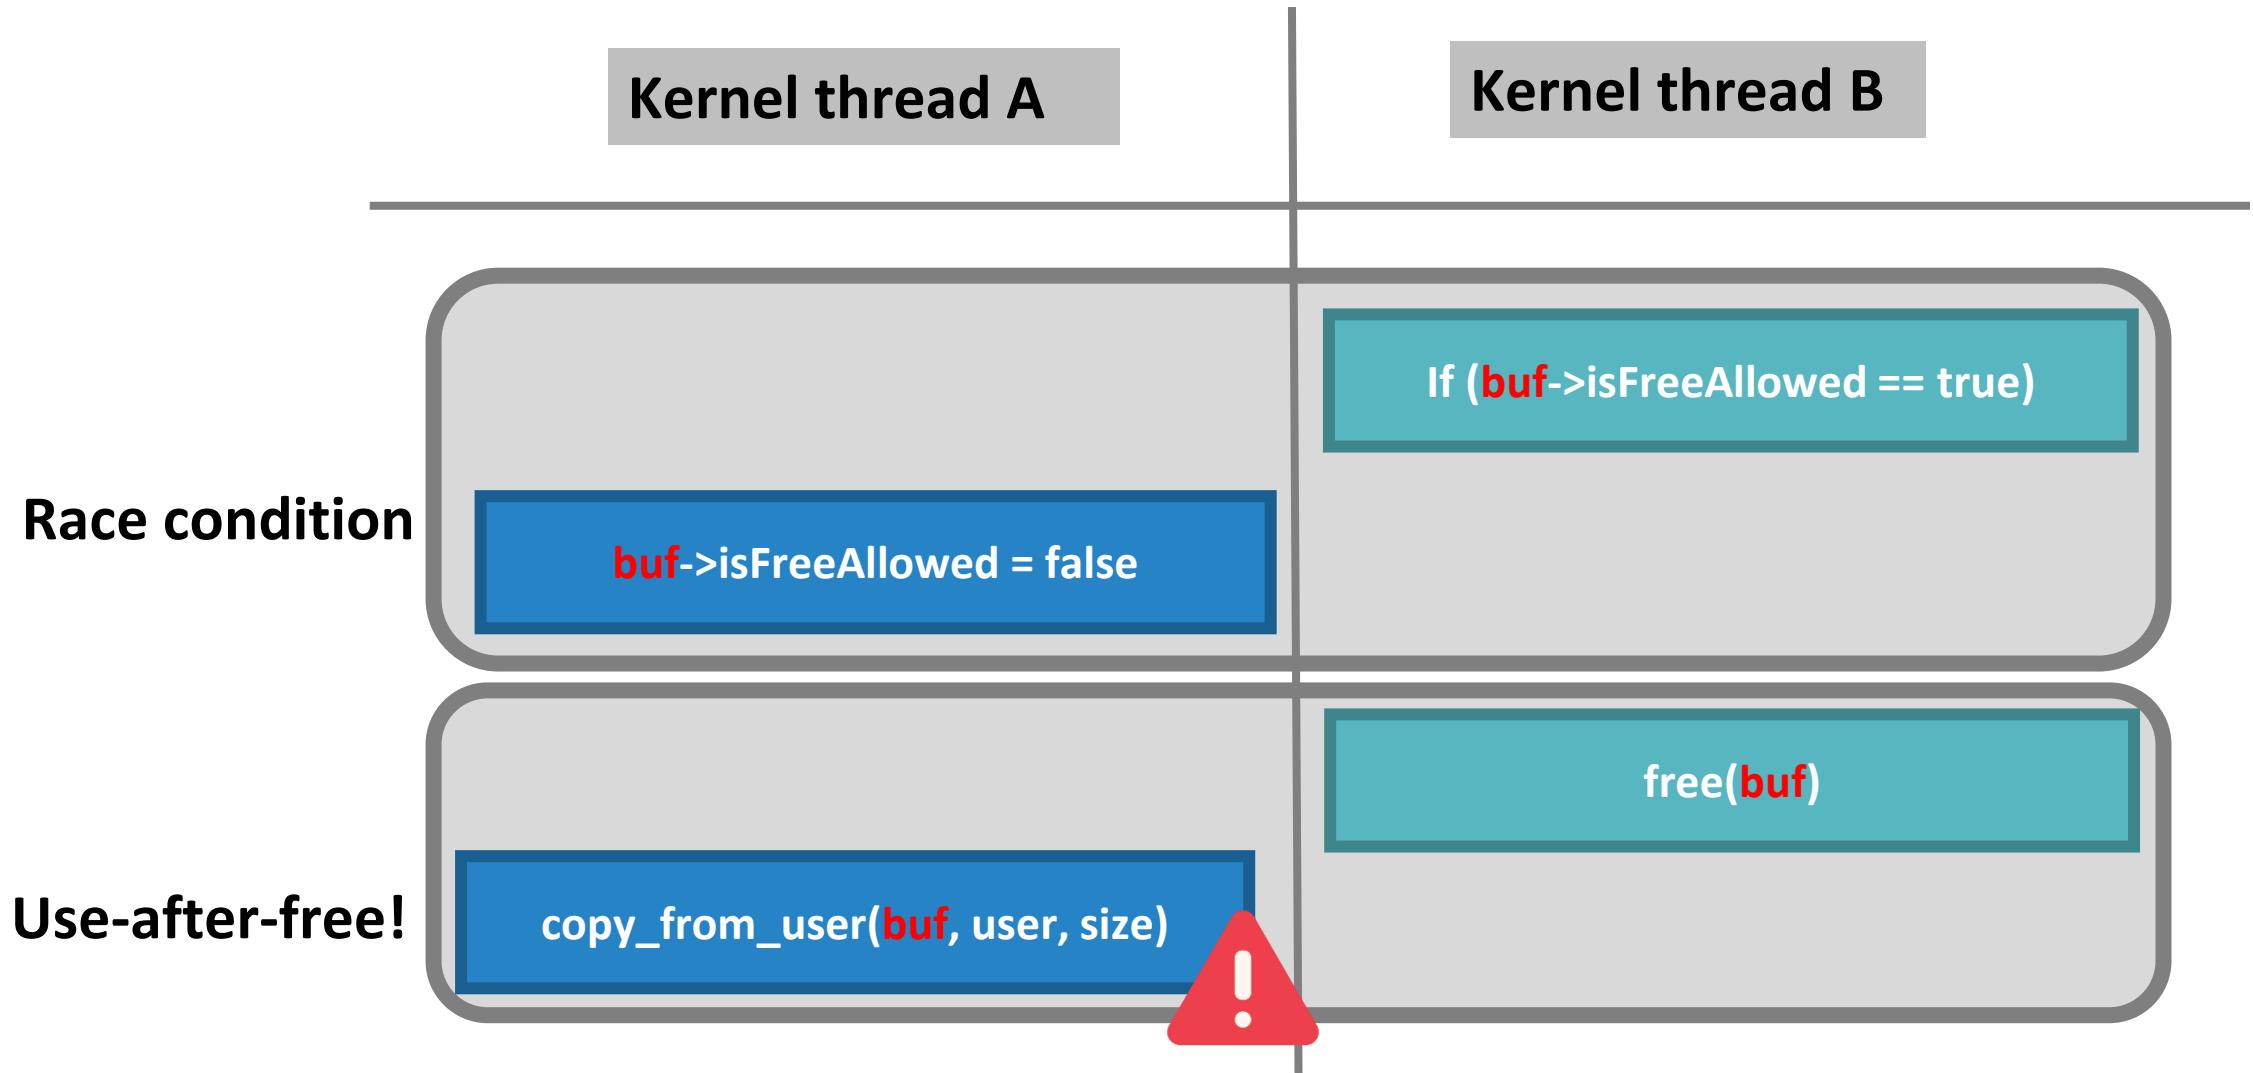
```
a = 1;
```

b will be 0

Race condition: CVE-2019-2025

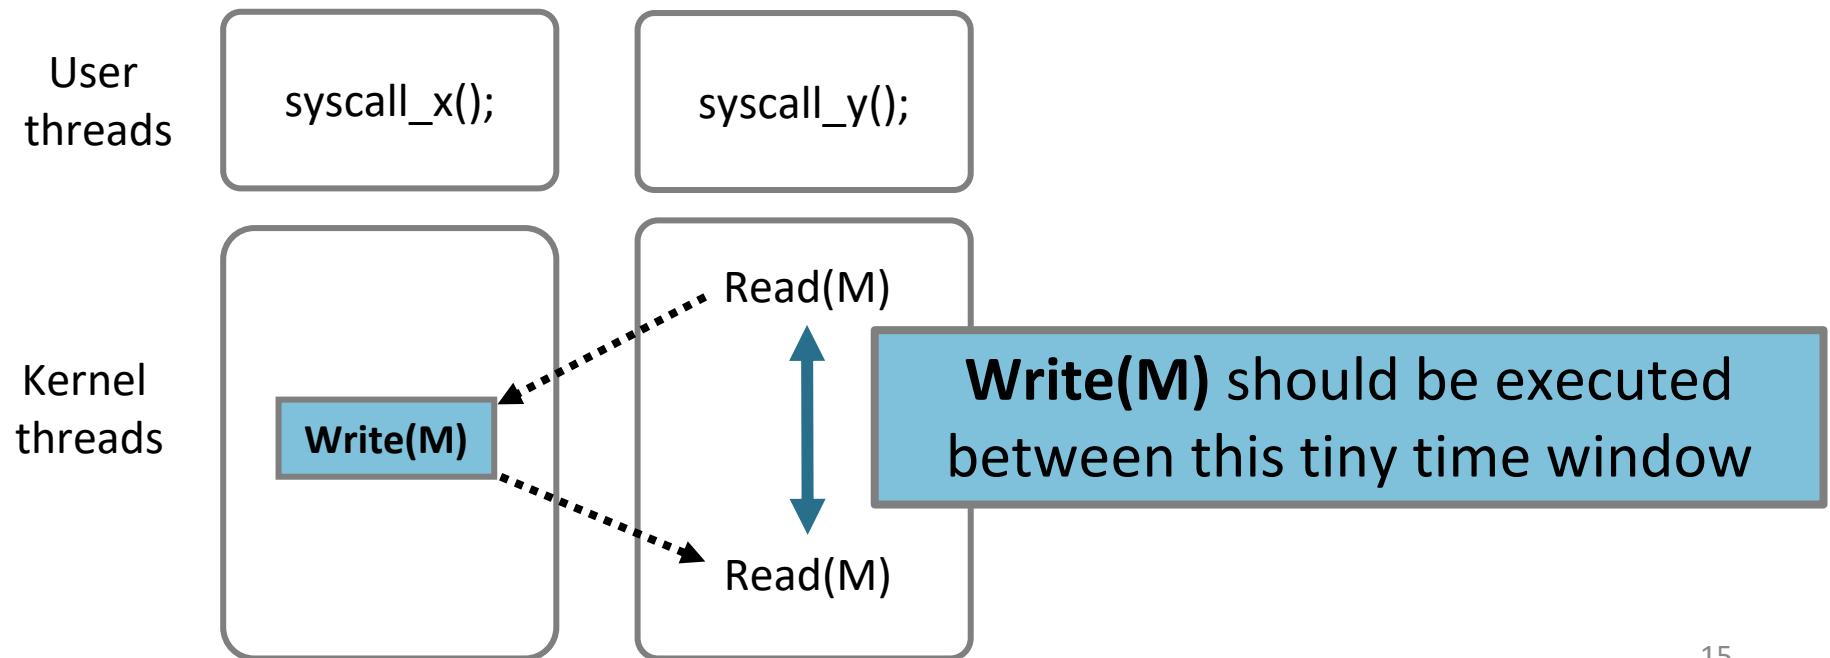


Race condition: CVE-2019-2025



Challenges: Kernel Races and Fuzzer

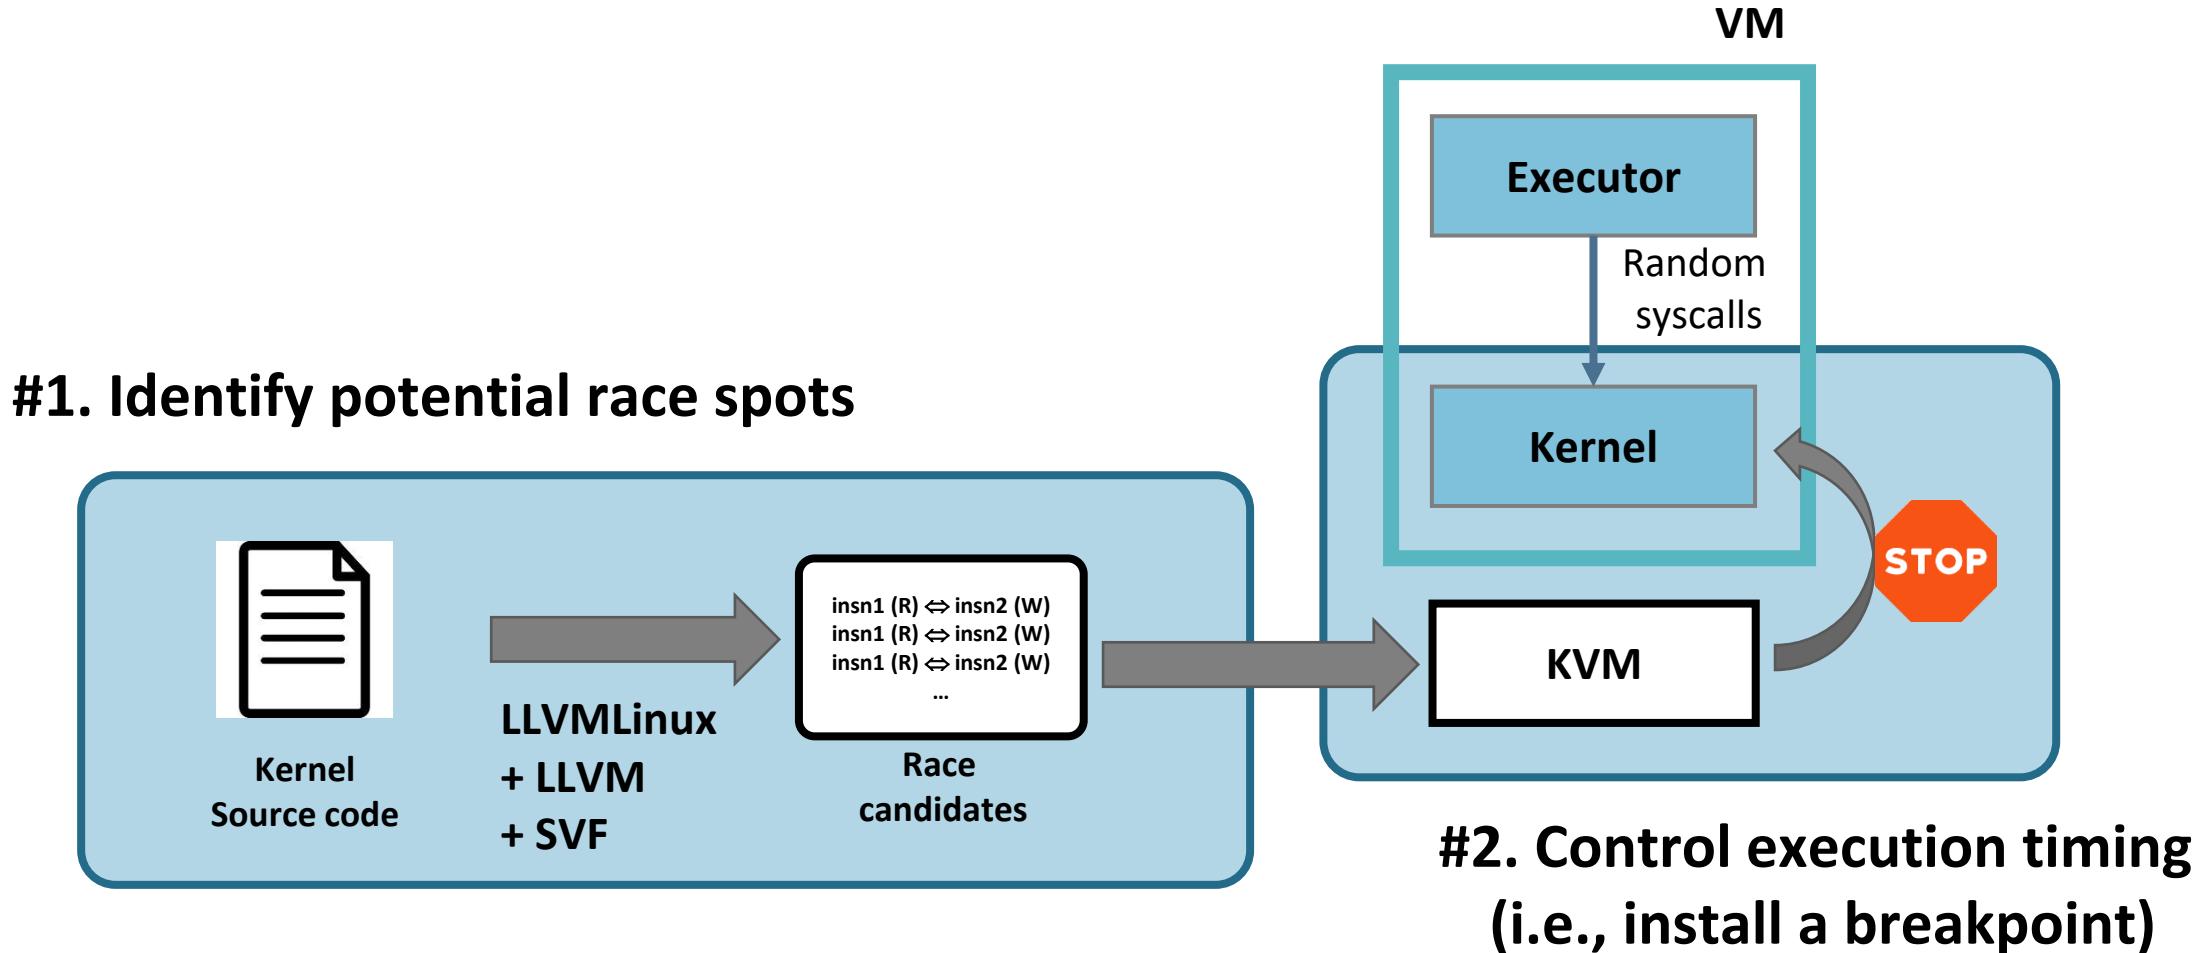
- Traditional fuzzers are not efficient to discover races
 - Execution timing (due to thread interleaving) matters to find races
 - Traditional fuzzers cannot take care of such timing



Our Approach: Razzer

- **Razzer: A new fuzzer to find RACE!**
 - #1. Static analysis to find potential race spots
 - Points-to analysis over an entire kernel
 - #2. For each potential race spots, setup per-core breakpoint
 - Deterministic kernel execution w.r.t. a given race spot
 - Run the kernel on modified hypervisor (QEMU/KVM)
 - #3. Check if a race truly occurs
 - Prioritize a truly racing pair to be further fuzzed later

Razzer: Kernel fuzzer for race conditions



Results

30 new races!
16 are fixed!

Use-after-free

Heap overflow

Double free

Kernel crash summary	Crash type	Kernel version	Kernel subsystem	Confirmed	Patch submitted	Fixed
KASAN: slab-out-of-bounds write in tty_insert_flip_string_flag	Use-After-Free	v4.8	drivers/tty/	✓	✓	✓
WARNING in __static_key_slow_dec	Reachable Warning	v4.8	net/	✓		
Kernel BUG at net/packet/af_packet.c:LINE!	Reachable Assertion	v4.16-rc3	net/packet/	✓	✓	✓
WARNING in refcount_dec	Reachable Warning	v4.16-rc3	net/packet/	✓	✓	✓
unable to handle kernel paging request in snd_seq_oss_readq_puts	Page Fault	v4.16	sound/core/seq/oss/	✓	✓	✓
KASAN: use-after-free Read in loopback_active_get	Use-After-Free	v4.16	sound/drivers/	✓	✓	✓
KASAN: null_ptr-deref Read in rds_ib_get_mr	Null ptr deref	v4.17-rc1	net/rdma/	✓ (assisted Syzkaller)	✓	✓
KASAN: null_ptr-deref Read in list_lru_del	Null ptr deref	v4.17-rc1	fs/			
BUG: unable to handle kernel NULL ptr dereference in corrupted	Null ptr deref	v4.17-rc1	net/sctp/			
KASAN: use-after-free Read in nd_jump_root	Use-After-Free	v4.17-rc1	fs/	✓	✓	✓
KASAN: use-after-free Read in link_path_walk	Use-After-Free	v4.17-rc1	fs/	✓	✓	✓
BUG: unable to handle kernel paging request in __inet_check_established	Page Fault	v4.17-rc1	net/ipv4/			
KASAN: null_ptr-deref Read in ata_pio_sector	Null ptr deref	v4.17-rc1	net/drivers/ata/			
WARNING in ip_recv_error	Reachable Warning	v4.17-rc1	net/	✓	✓	✓
WARNING in remove_proc_entry	Reachable Warning	v4.17-rc1	net/sunrpc/			
KASAN: null_ptr-deref Read in ip6gre_exit_batch_net	Null ptr deref	v4.17-rc1	net/ipv6/			
KASAN: slab-out-of-bounds Write in __register_sysctl_table	Heap overflow	v4.17-rc1	net/ipv6/			
KASAN: use-after-free Write in skb_release_data	Use-After-Free	v4.17-rc1	net/core/			
KASAN: invalid-free in ptlock_free	Double free	v4.17-rc1	mm/			
Kernel BUG at lib/list_debug.c:LINE!	Reachable Assertion	v4.17-rc1	drivers/infiniband/			
INFO: trying to register non-static key in __handle_mm_fault	Reachable INFO	v4.17-rc1	mm/			
KASAN: use-after-free Read in vhost-chr_write_iter	Use-After-Free	v4.17-rc1	drivers/vhost/	✓	✓	✓
BUG: soft lockup in vmemdup_user	Soft lockup	v4.17-rc1	net/			
KASAN: use-after-free Read in rds_tcp_accept_one	Use-After-Free	v4.17-rc1	net/rds/			
WARNING in sg_rq_end_io	Reachable Warning	v4.17-rc1	drivers/scsi/			
BUG: soft lockup in snd_virmidi_output_trigger	Soft lockup	v4.18-rc3	sound/core/seq/	✓ (assisted Syzkaller)	✓	✓
KASAN: null_ptr-deref Read in smc_ioctl	Null ptr deref	v4.18-rc3	net/smci/	✓	✓	✓
KASAN: null_ptr-deref Write in binderf_update_page_range	Null ptr deref	v4.18-rc3	drivers/android/	✓	✓	
WARNING in port_delete	Reachable Warning	v4.18-rc3	sound/core/seq/	✓ (assisted Syzkaller)		
KASAN: null_ptr-deref in inode_permission	Null ptr def	v4.18-rc3	fs/	✓	✓	✓

Razzer is open-sourced!



compsec-snu / razzer



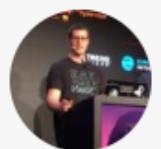
Fork

61



Starred

368



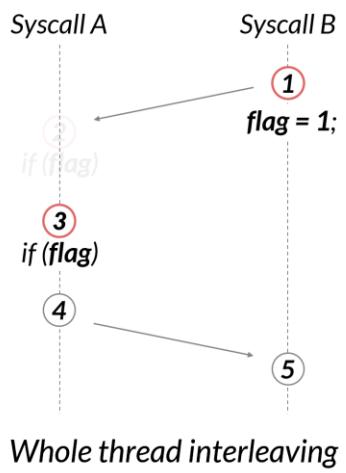
Yoav Alon @yoavalon · May 28, 2019



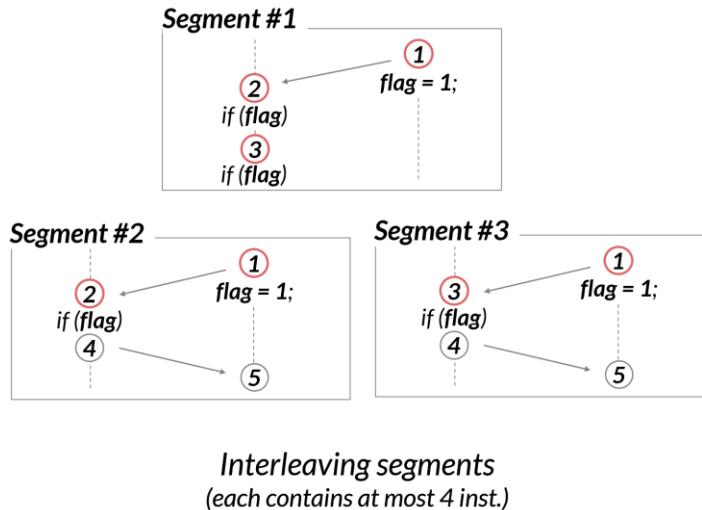
...

The source code for **Razzer** - race condition **fuzzer** for the linux kernel was released. It's one of the most innovative fuzzers I've seen in recent years. It uses a clever static analysis to gather race candidates and a fork of syzkaller & hv to test them.

Our follow-up work: SegFuzz [SP 23]



Whole thread interleaving



Interleaving segments
(each contains at most 4 inst.)

general protection fault in vmci_host_poll
KASAN: use-after-free Read in cfusbl_device_notify
KASAN: use-after-free Read in slcan_receive_buf
general protection fault in ctimeout_net_exit
KASAN: use-after-free Read in raw_notifier_call_chain
INFO: task hung in blk_trace_remove
INFO: task hung in blk_trace_setup
kernel BUG in pfkey_send_acquire
general protection fault in add_wait_queue_exclusive
KASAN: use-after-free Read in slip_ioctl
general protection fault in add_wait_queue
WARNING in isotp_tx_timer_handler
KASAN: use-after-free Read in snd_pcm_plug_read_transfer
Kernel BUG in find_lock_entries
KASAN: use-after-free Read in tcp_write_timer_handler
KASAN: use-after-free Read in event_sched_out
general protection fault in soft_cursor
KASAN: use-after-free Read in perf_event_groups_insert
BUG: unable to handle kernel paging request in usb_start_wait_urb
BUG: unable to handle kernel paging request in __kernfs_new_node
general protection fault in raw_seq_start

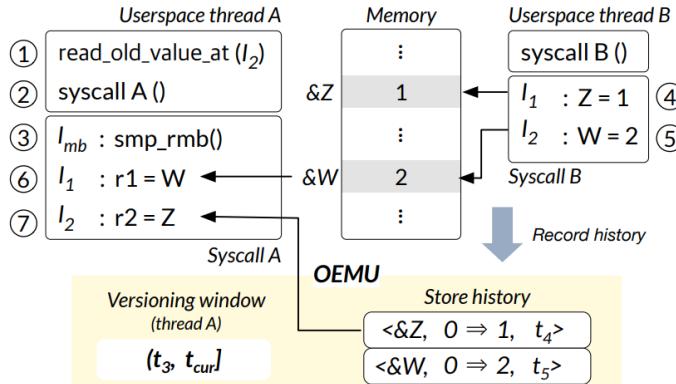
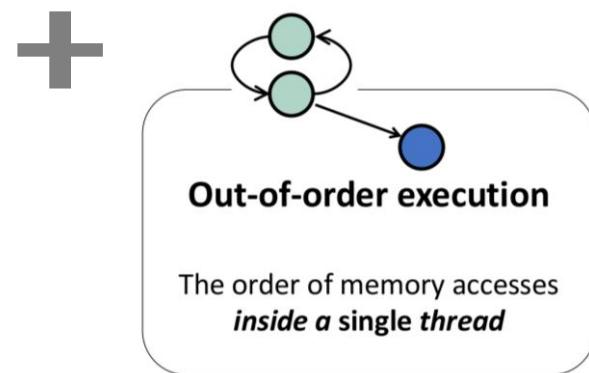
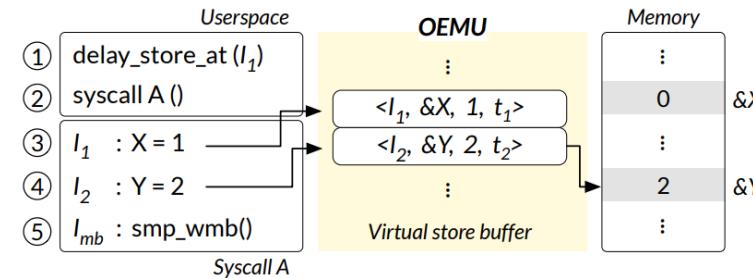
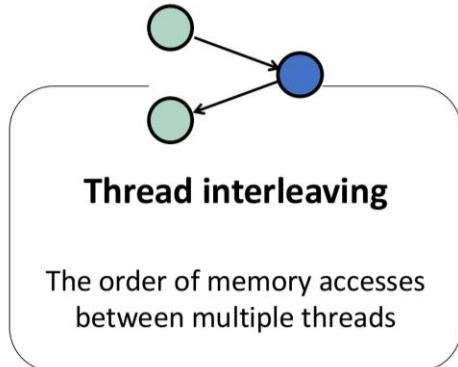
Developed new thread interleaving coverage
→ Efficiently exploring unexplored thread interleavings

21 new race vulnerabilities!

is presented to

Dae R. Jeong, Yewon Choi, Byoungyoung Lee,
Insik Shin, Youngjin Kwon

Our follow-up work: OZZ [SOSP 24]



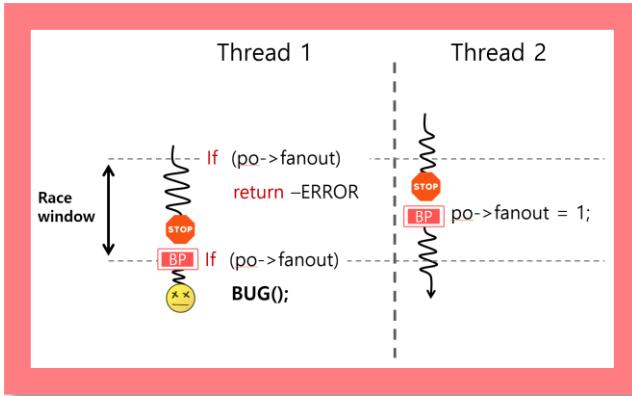
Challenge:

- Complex non-deterministic behaviors when identifying out-of-order bugs

Solution:

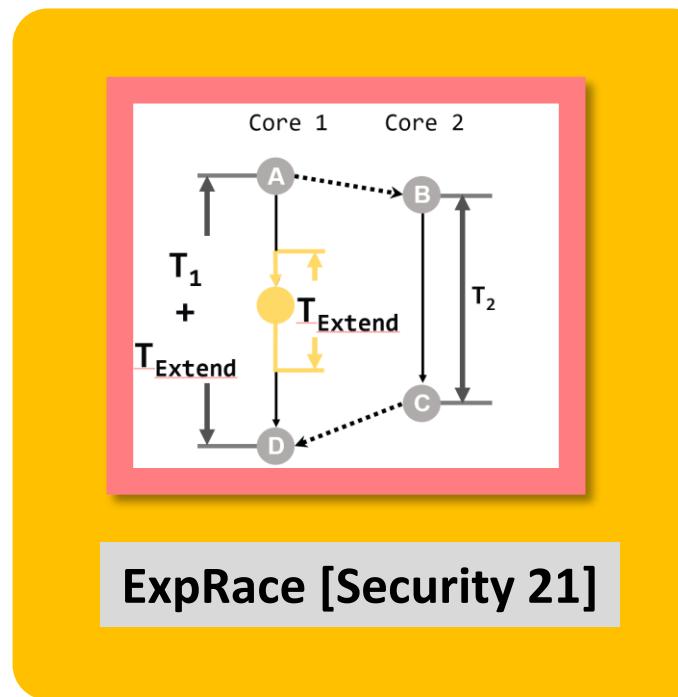
- Emulate load/store to be deterministic
- Fuzzing with SegFuzz [SP 23]

Our Research on Linux Kernel Security



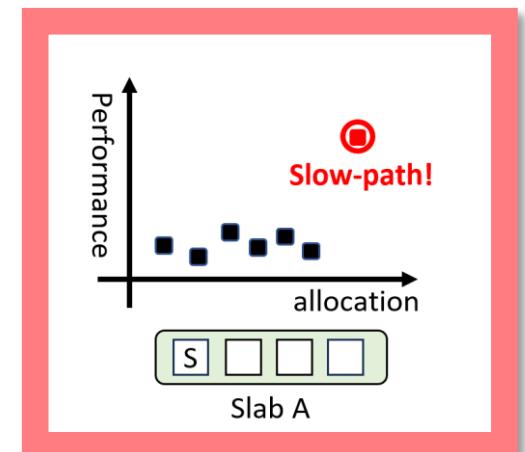
Razzer [SP 19]
SegFuzz [SP 23]
OZZ [SOSP 24]

Fuzzing to find race vulnerabilities



ExpRace [Security 21]

Exploiting race vulnerabilities



Pspray [Security 23]

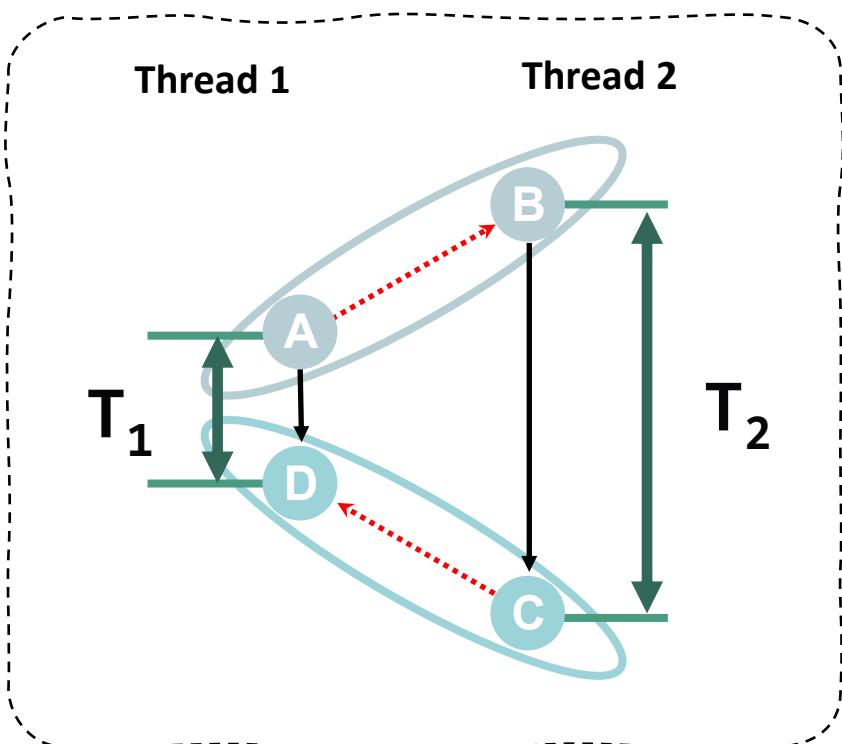
Problem : Multi-Variable Race

Race CVE	Does brute-force work?
CVE-2019-11486	✓
CVE-2017-7533	✓
CVE-2017-2636	✓
CVE-2016-8655	✓
CVE-2019-6974	✗
CVE-2019-2025	✗
CVE-2019-1999	✗
CVE-2017-15265	✗

• Some races are exploitable with brute-forcing

• However, some others are NOT.

Problem : Multi-Variable Race



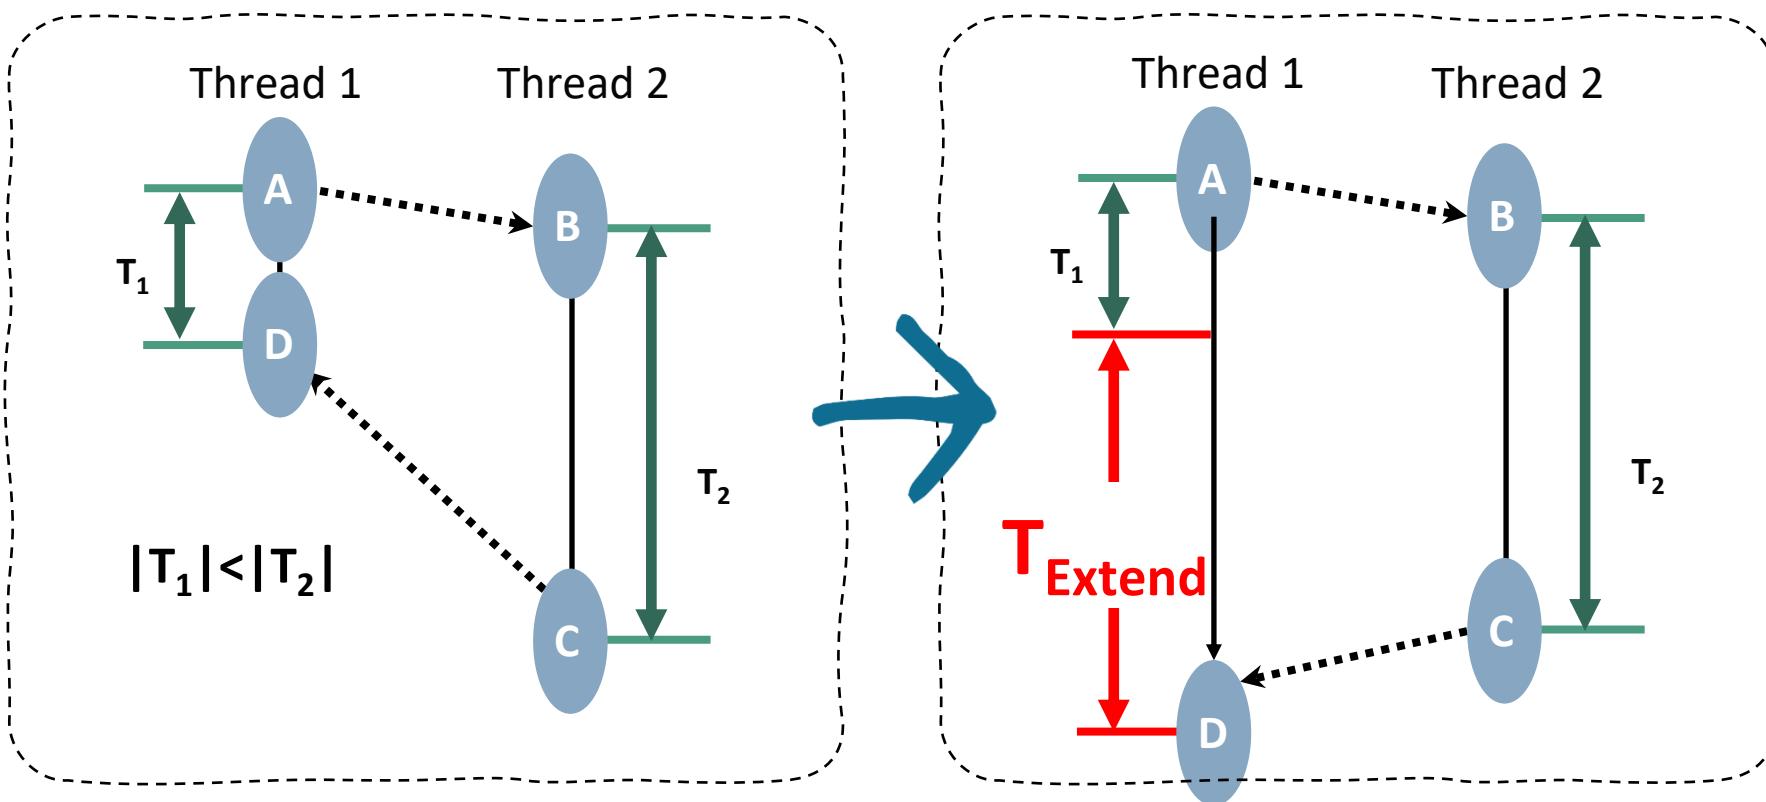
Access to memory M_1



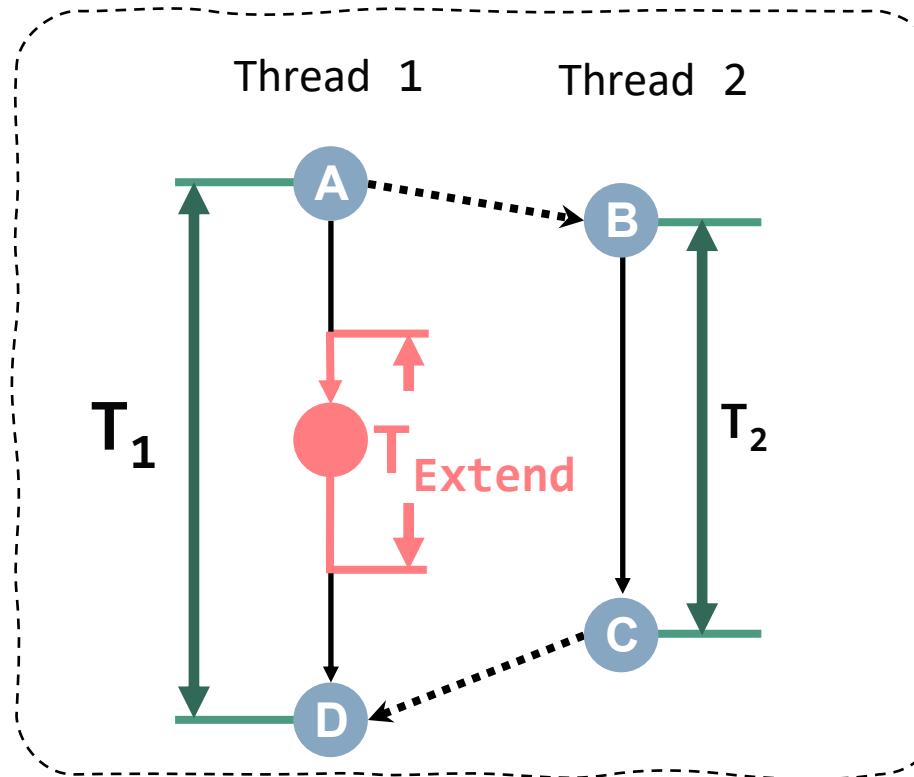
Access to memory M_2

- Some races are (nearly) impossible to exploit
 - If $T_1 < T_2$

Goal: Make Exploit Work Again!

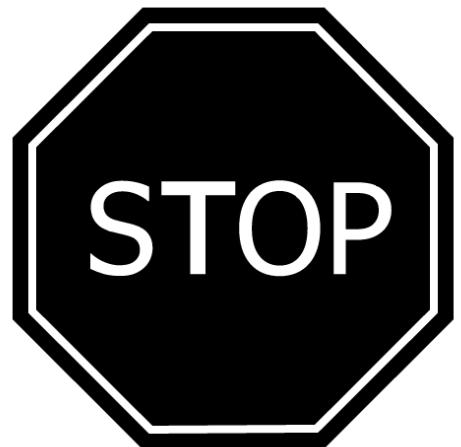


Idea: How to extend the time window?

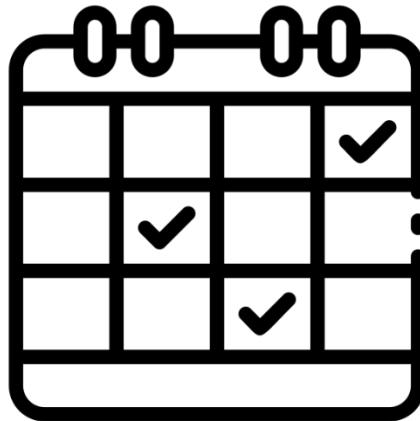


Preempting the thread!

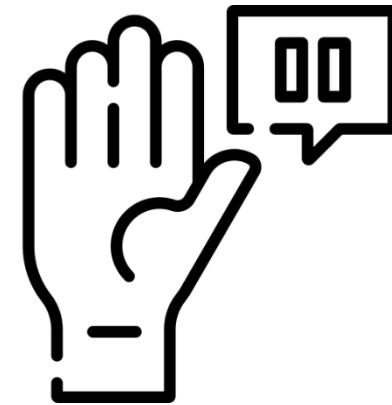
Background: Kernel Preemption



Kernel breakpoint



Kernel thread
schedule



Interrupt

Challenge: Users cannot preempt!



Kernel breakpoint



Kernel thread
schedule



Interrupt

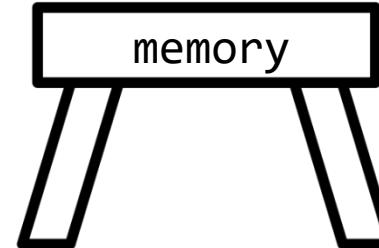
Back to OS basics:
Users cannot control kernel's preemption

Idea again: User-controlled Preemption through Interrupts

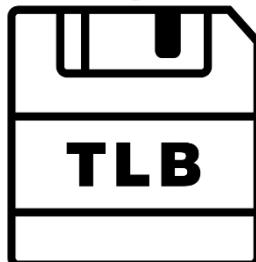
We found four methods to **indirectly raise interrupts through syscalls**



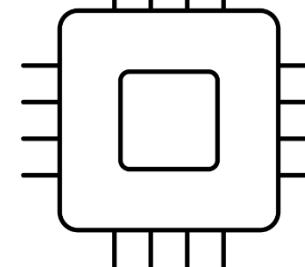
- Reschedule interrupt



- Membarrier interrupt

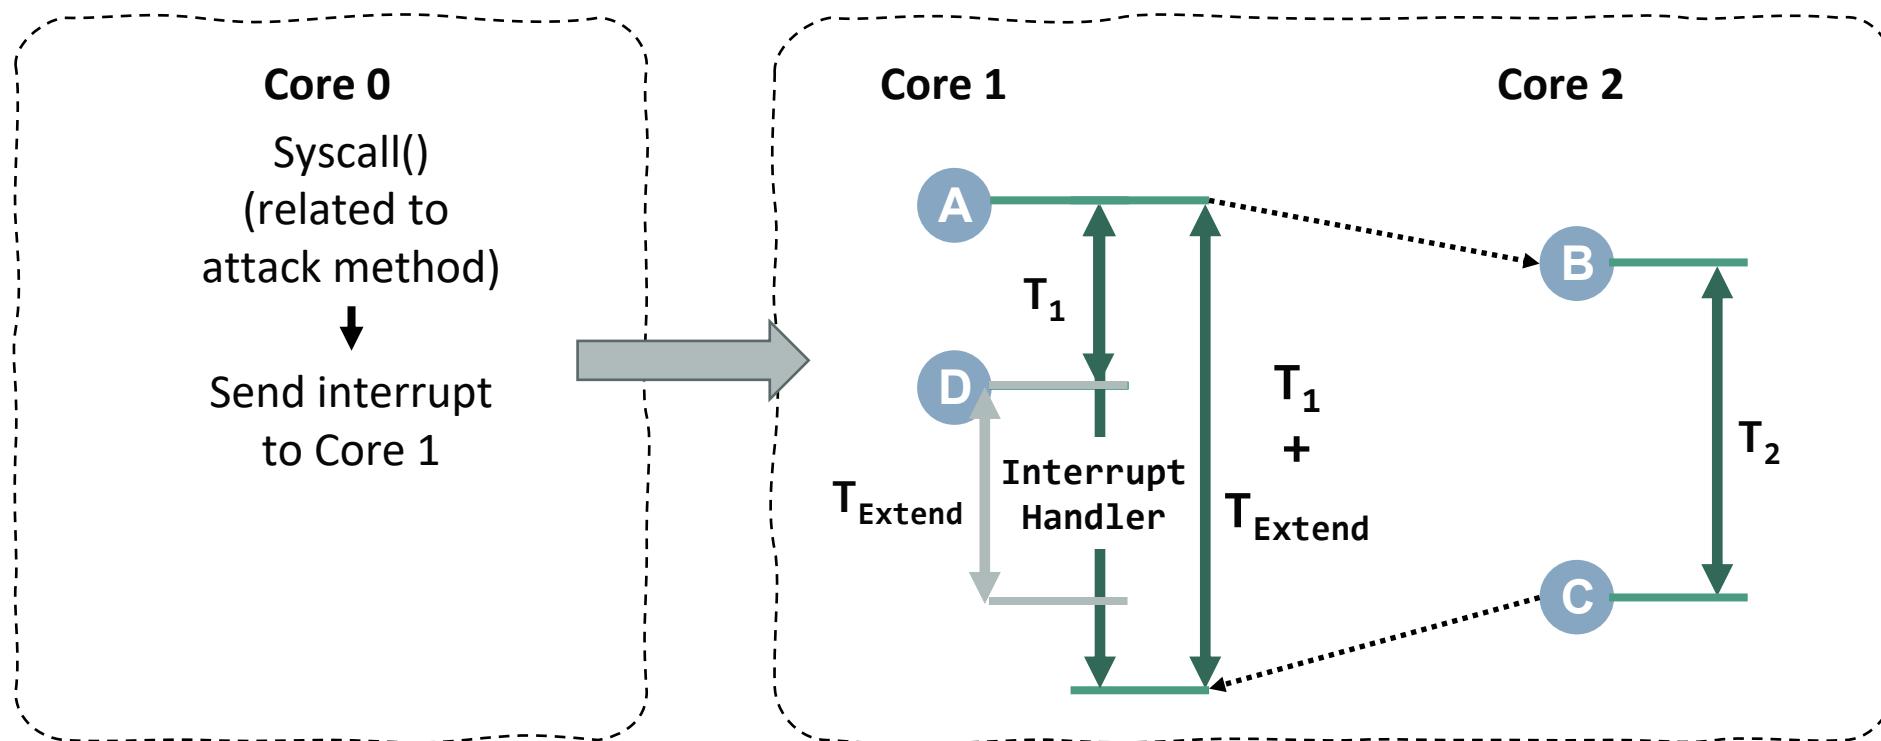


- TLB shootdown interrupt



- Hardware interrupt

Our Approach : Interrupt



Enlarging race window through syscalls and its follow-up interrupts

Real-World Races in Linux

Vulnerability	Baseline	Reschedule	membarrier	TLB shootdown	HW interrupt
CVE-2019-6974	✗	✗	✗	✗	✓
CVE-2019-2025	✗	✓	✓	✓	✓
CVE-2019-1999	✗	✗	✗	✓	✓
CVE-2017-15265	✗	✓	✓	✓	✓
11eb85ec...	✗	✗	✗	✓	✓
1a6084f8...	✗	✗	✗	✓	✓
20f2e4c2...	✗	✗	✗	✓	✓
484298f...	✗	✓	✓	✓	✓
da1b9564...	✗	✗	✗	✗	✓
e20a2e9c...	✗	✗	✗	✓	✓

✗ denotes exploitation has failed for given 24 hours

All 10 cases were successfully exploited with ExpRace!

After the publication of ExpRace

 **Theori** 
@theori_io

 New Linux Kernel vulnerability (CVE-2024-27394) discovered & patched by Theori!

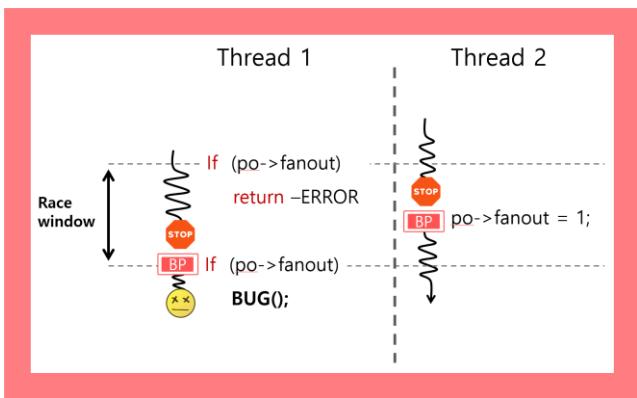
 blog.theori.io/deep-dive-into...

Our researcher [@v4bel](#) at #Theori identified a critical #UAF vulnerability in TCP-AO caused by a race condition in the #RCU API. Using techniques from the ExpRace paper, we extended the race window to demonstrate its exploitability.

Curious how we did it?
Read our deep dive for the full details!

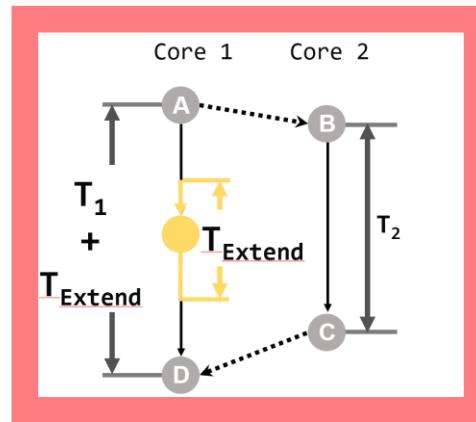
#Theori #Cybersecurity #LinuxKernel #TCP #VulnerabilityResearch #CVE #TechBlog

Our Research on Linux Kernel Security



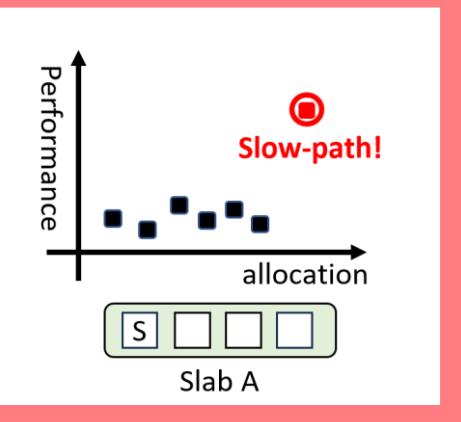
Razzer [SP 19]
SegFuzz [SP 23]
OZZ [SOSP 24]

Fuzzing to find race vulnerabilities



ExpRace [Security 21]

Exploiting race vulnerabilities



Pspray [Security 23]

Summary of PSpray

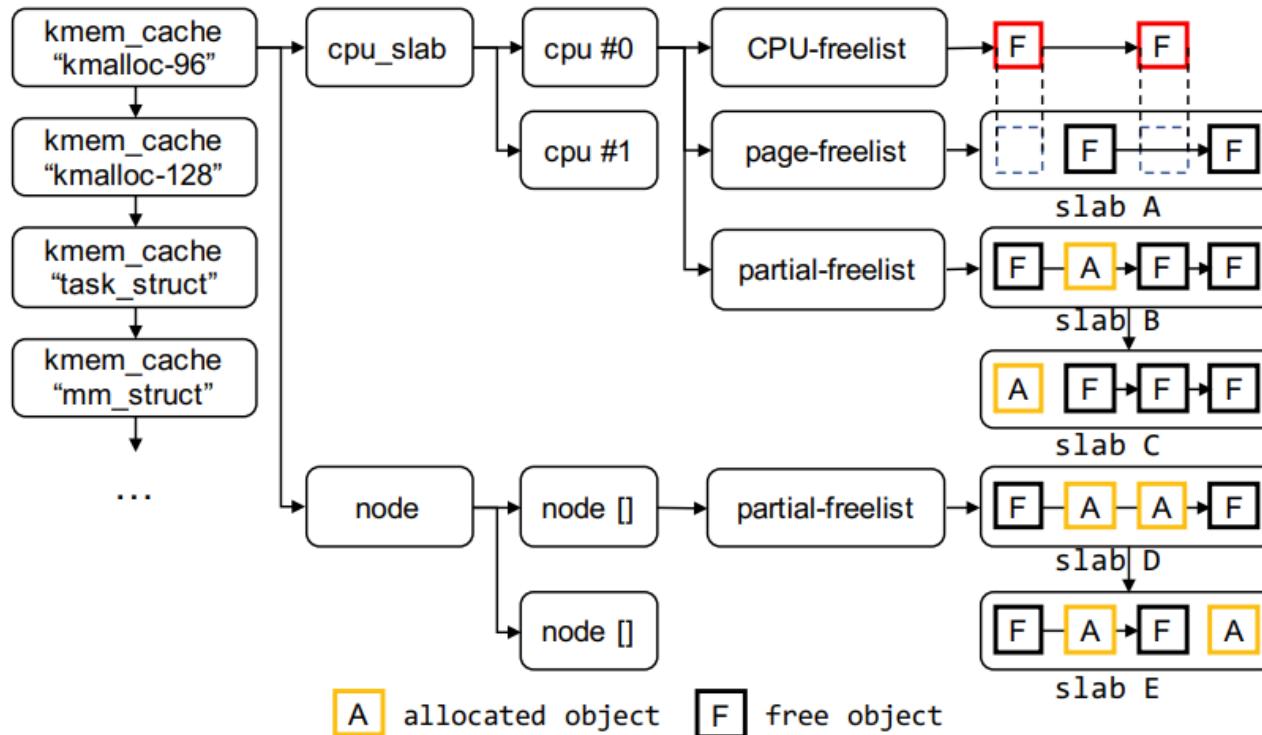
Motivation

- **Reliable kernel exploitation** is important for attackers.
- Reliable use-after-free exploitation is challenging.
- Why? **Attackers do not know the status of kernel heap.**

Our finding

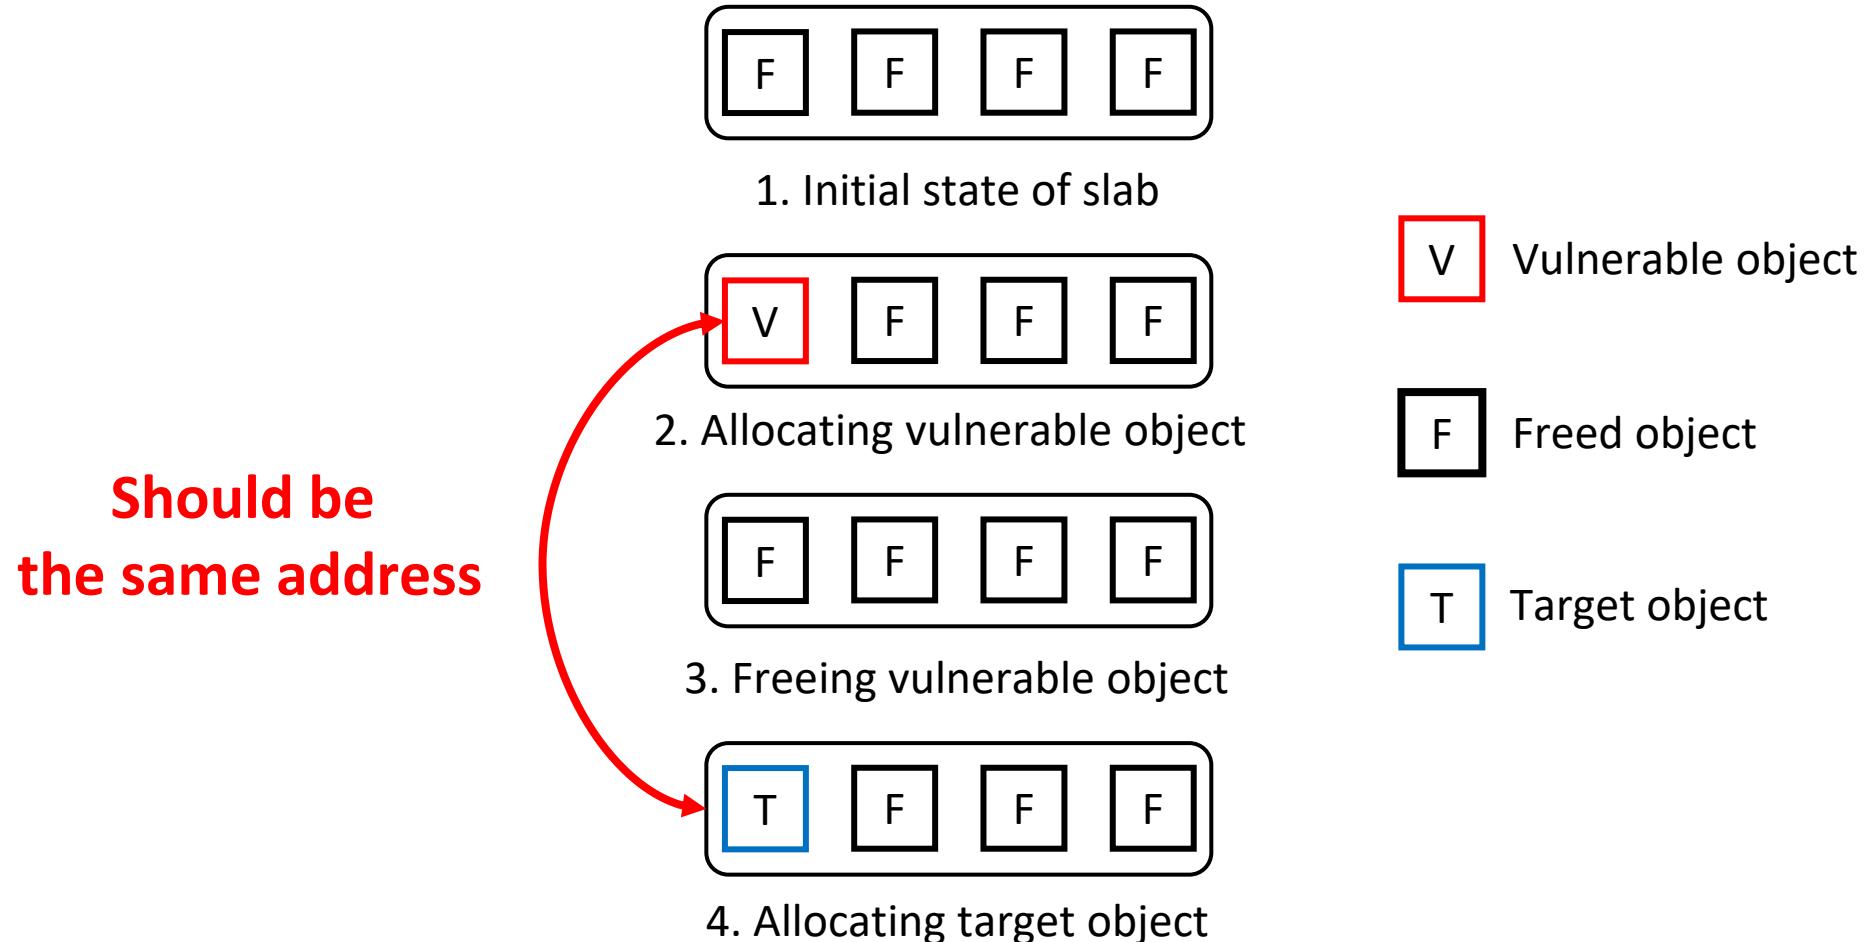
- **Timing side-channel attacks** to **learn the status of kernel heap.**
- Attackers can launch reliable kernel exploitation.

Slab: Kernel's memory allocator



Slab's allocation process is a bit complicated 😞

Challenges in exploiting **use-after-free**



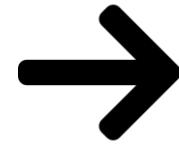
The **vuln** and **target** objects must be placed at the **same address** to exploit.

Idea: Timing Side-Channel

- Kernel implements **multiple allocation paths**:
 - Fast, medium, slow paths depending on the heap status
- We found **timing side-channels to identify allocation path!**
 - By measuring the time taken to allocate,
we know which allocation path has been taken.
- This allows to place vulnerable and target objects **at the same address!**

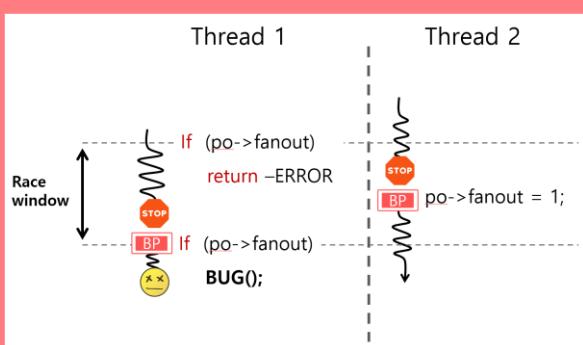
Exploit success rate after Pspray

Vulns	Type	# of alloc	Baseline (Success rate)	Pspray (Success rate)
CVE-2019-2215	UAF	2	93.28%	100.0%
CVE-2018-6555	UAF	13	63.50%	99.94%
83bec2...	UAF	8	13.70%	98.16%
77e2cf...	UAF	1	95.74%	100.0%
CVE-2017-6074	DF	4	80.64%	100.0%
6b8d6b...	DF	1	96.28%	99.98%



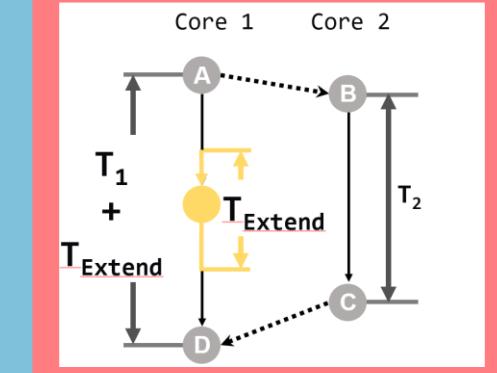
Conclusion

- Linux kernel security is important!

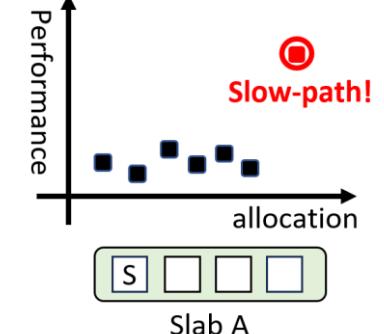


Razzer [SP 19]
SegFuzz [SP 23]
OZZ [SOSP 24]

Fuzzing to find race vulnerabilities



ExpRace [Security 21]



Pspray [Security 23]

Exploiting use-after-free

감사합니다

Byoungyoung Lee
Seoul National University
byoungyoung@snu.ac.kr
<https://compsec.snu.ac.kr>

